BatchGuard

Make the next inspection count.

INSPECTOR DECISION SUPPORT

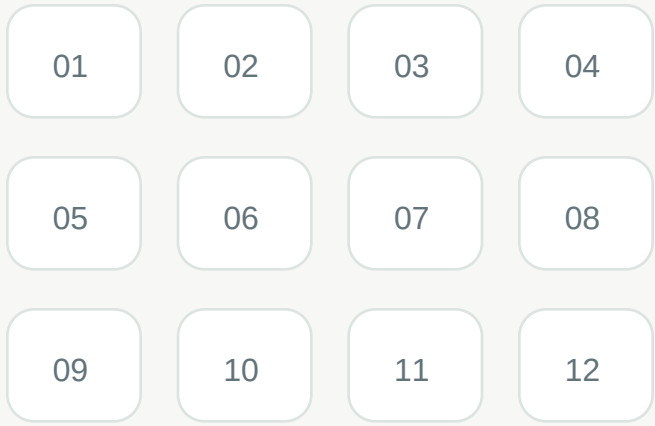

12 batches.
4 inspection slots.

Which four—and why?

INSPECTION CAPACITY

NEXT NEXT NEXT NEXT

Illustrative capacity scenario · human review remains in control

Make evidence easier to act on.

A clearer handoff from a shortlist to a documented human review.

01

Drug inspectors

See what to review first, why it ranked, and which source report supports the case.

02

Review teams

Trace recommendation context and keep reviewer actions attached to the case.

03

Lab coordinators

Receive the report and scan context behind a proposed sample.

EXPECTED VALUE

Prioritize with evidence. Explain the choice. Record the decision.

Intended workflow value; time savings, clinical outcomes and adoption are not yet measured.

Real reports. Simulated scans.

The two evidence streams stay separate and visibly labelled.

2,551
analytical records

14 months
Jul 2025 – Aug 2026

2,451
match captured official fields

REAL · CDSCO REPORTS

Product, batch, maker text

Reported quality findings

Workbook, sheet and row references

100 records flagged for review

SIMULATED · PACK SCANS

Pack IDs, readers, times

Scripted location conflicts and bursts

Used to exercise pipeline behavior

Not field accuracy or counterfeit proof

Report recurrence is not a medicine-safety label. Maker text is not verified legal identity.

An inspection queue you can explain.

History model ranks report recurrence; explicit rules flag unusual scan patterns.

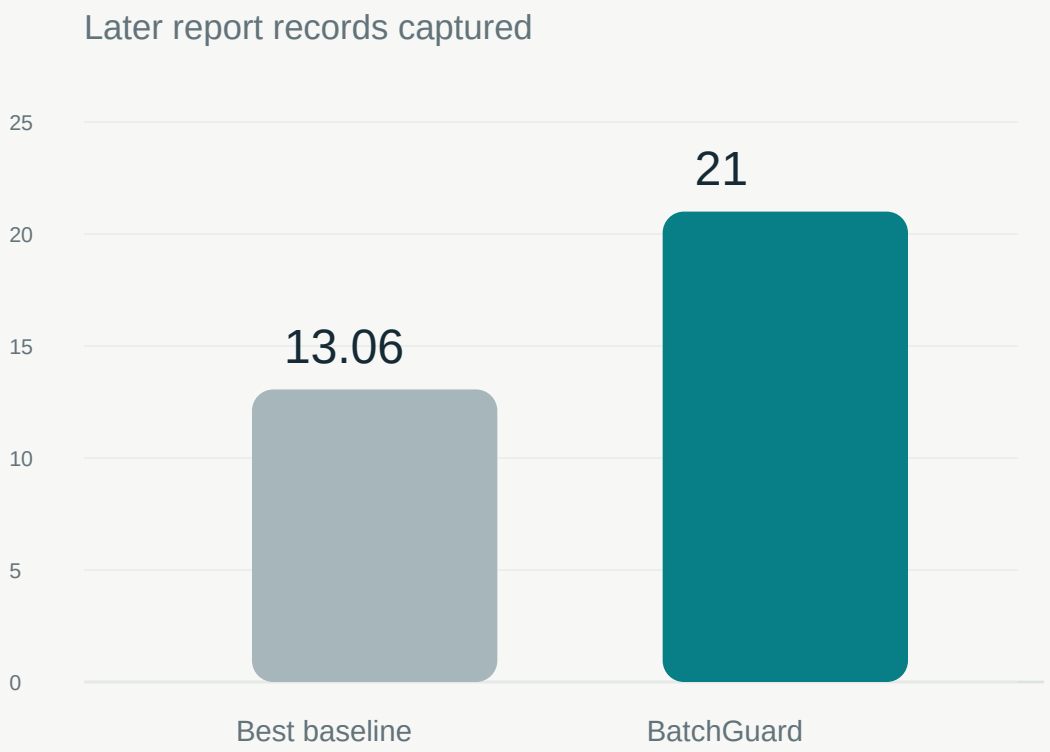

+61%

more later report records captured at the same ranking budget

BUDGET: 25 MAKER TEXTS / MONTH

History-only test, July–August 2026. Baseline: recent report counts. This compares report recurrence, not unsafe batches found. Separate from the simulated scan demo.

One input. One traceable review.

A simulated case shows how a new signal changes the queue.

1 CSV input → 2 Validate → 3 Score → 4 Rank → 5 Open evidence → 6 Record review

SCRIPTED DEMO · CSV INPUT

Rank 12 → Rank 4

Inside four available slots

WHAT CHANGES THE RANK?

A simulated same-pack location conflict triggers a review rule. Removing the conflict makes the scripted check fail.

The demo exercises the pipeline; it does not validate counterfeit detection in the field.

Executed on IBM LinuxONE.

A real s390x deployment run—not a hardware-acceleration claim.

IBM LINUXONE · s390x

Models retrained

on the host's scikit-learn 1.9.0

Preflight passed

Service health returned HTTP 200

What LinuxONE does here

- Hosts the central Python service
- Runs model inference and queue ranking
- Keeps review state with the service

Standard CPU · NNPA not used or measured

The model is portable; it does not require IBM Z. The demonstrated value is running the shared decision service on the LinuxONE host.

A decision that can be inspected.

Sources and limits are part of the product.

PUBLIC DATA	CDSCO NSQ and spurious-drug reports; original source references retained.	https://cdscoonline.gov.in
EVALUATION	Chronological report-recurrence test, compared with a recent-count baseline.	July–August 2026 · top 25 / month
PLATFORM	LinuxONE s390x run verified; no NNPA or hardware speedup claimed.	IBM LinuxONE · IBM Snap ML
MAP	Natural Earth 1:10m India point-of-view geometry; contextual map only.	naturalearthdata.com/downloads/10m-cultural-vectors/

BatchGuard prioritizes review. Inspectors and laboratories determine what happens next.